

- Introduction
 - What
 - Common Features
 - Provides
 - Why
 - Supports
 - Types of NoSQL
 - 1. Key-Value Stores
 - Example
 - 2. Column Family / Columnar Databases
 - Use Cases
 - 3. Graph Databases
 - Use Cases
 - 4. Document Store Databases
 - Example
- MongoDB
 - The Document Model and MongoDB
 - Structure and Syntax
 - Collections
 - Document
 - Collection
 - MQL: MongoDB Query Language
 - MQL Create
 - MQL Find
 - Query filter document
 - Specifying query operators
 - Comparison Operators
 - Logical Operators
 - Examples
 - MQL Update
 - MQL Delete
 - MQL Drop
- MongoDB Aggregation Expressions
 - Syntax
 - Suggested ordering of stages in the pipeline
 - Stages
 - \$match
 - \$expr
 - \$project
 - \$limit
 - \$group
 - \$addFields
 - \$function
 - \$unwind
 - Options
 - Examples
 - \$lookup

- [Index](#)
 - [Aggregation Expression: Operators](#)

Introduction

What

- NoSQL stands for **not only SQL**
- family of databases that vary widely in style and technology

Common Features

1. Non relational in featured
2. not RDBMS

Provides

1. New ways of storing and querying data
2. Databases do not require fixed schemas
3. Horizontally scale easily
4. Provide native fault tolerance and availability in distributed systems

Why

Applications shifting from internal employees to the public internet users requires high availability and performance

Supports

- flexible data model
- built in horizontal and vertical scaling
- developer productivity
- distributed environments
- **Polyglot persistence**: mixing RDBMS solutions with NoSQL solutions

Types of NoSQL

Key-Value, Document, Column, Graph

1. Key-Value Stores

Simplest NoSQL databases - in which it stores data as a collection of key-value pairs

1. Keys are unique
2. Directly points to its associated data

Example

Redis or Memcached: caching frequently used data

DynamoDB: E-commerce platforms, gaming applications, high scalable

2. Column Family / Columnar Databases

1. Organize data in columns rather than rows
2. Efficient for handling data sets with dynamic schemas

Use Cases

1. **Apache Cassandra**: IoT applications
2. **Hbase**: applications that store and analyze user preferences and behaviours
3. Large-scale data analysis

3. Graph Databases

Designed to manage highly interconnected data, representing relationships

Use Cases

1. Social Networking
2. Recommendation systems

4. Document Store Databases

1. Store data in document format (JSON, BSON, or XML)
2. Each document contains key-value pairs or key-document pairs
3. These databases are schema-less, allowing flexibility in data structures within a collection

Example

1. MongoDB
2. CouchBase
3. Amazon DocumentDB

MongoDB

The Document Model and MongoDB

Structure and Syntax

1. JSON Notation
2. Field-value pairs are separated by colon
3. Fields must be enclosed within quotation marks
4. Two field-values are separated by commas

Collections

Document

A way to organize and store data as a set of field-value pairs in MongoDB

Collection

An organized store of documents in MongoDB, usually with common fields between documents

MQL: MongoDB Query Language

MQL Create

`insertOne()`: Insert one document into a collection

```
db.cows.insertOne({
  name: "daisy",
  milk: 8
});
```

`insertMany()`: Insert an array of documents into a collection

```

db.people.insertMany([
  {
    "user_id": "Eoin",
    "age": 29,
    "Status": "A"
  },
  {
    "user_id": "Daniel",
    "age": 25,
    "Status": "A",
    "Country": "USA"
  }
], ordered: false);

```

writeConcern: Sets the level of acknowledgement requested from MongoDB for write operations

ordered: For insertMany(), there is an additional option for controlling whether the documents are inserted in ordered or unordered fashion

MQL Find

Query filter document

```
db.<collection>.find({<field1>:<value>,...})
```

Specifying query operators

```
db.<collection>.find({ <field1>: { <operator1>: <value1> }, ... })
```

Comparison Operators

```
$eq, $gt, $lt, $gte, $lte, $ne, $in, $nin
```

Logical Operators

```
$and, $not, $or, $nor
```

Examples

Using \$and

```

// Find cows that have milk between 6 and 9
cowCol.find(
  {
    $and: [
      {milk: {$gt: 6}},
      {milk: {$lt: 9}}
    ]
  }
)

```

Using implicit AND and OR

```

// Find an inventory that has status "A" and either has a quantity of less than 30 or the item name starts with a p
dn.inventory.find({
  status: "A",
  $or: [
    {price: {$lt: 30}},
    {item: /^p/}
  ]
})

```

Using explicit AND and OR

```
// Find an inventory that has (quantity is less than 10 or more than 50) and (status is A or size.w is less than 15)
dn.inventory.find({
  $and: [
    { $or: [
      { qty: { $lt: 10 } },
      { qty: { $gt: 50 } }
    ] },
    { $or: [
      { status: "A" },
      { "size.w": { $lt: 15 } }
    ] }
  ]
})
```

Using the logical NOT operator

```
// Find items that do not start with e letter p
db.inventory.find({
  item: { $not: { $regex: /^p.*/ } }
})
```

MQL Update

`updateOne()`: Update one document into a collection

```
// Set and change the milk number
db.cows.updateOne(
  {
    name: "daisy",
    milk: 12
  },
  {
    $set: { milk: 8 }
  }
);
```

`updateMany()`: Update an array of documents into a collection

```
// Increment all milk count in the collection
db.cows.updateMany(
  {},
  { $inc: { milk, 1 } }
);
```

MQL Delete

`deleteOne()`: Delete one document into a collection

```
// Set and change the milk number
db.cows.deleteOne(
  {
    name: "daisy",
    milk: 9
  }
);
```

`deleteMany()`: Delete an array of documents into a collection

```
// Increment all milk count in the collection
db.cows.deleteMany({});
```

MQL Drop

`drop()`: Removes the entire database object

```
db.some_collection.drop();
```

MongoDB Aggregation Expressions

The `$expr` operator allows the use of aggregation expressions in MQL

```
{ $expr: { <expression> } }
```

Like functions and take arguments, typically these are an array of arguments.

Syntax

An aggregation pipeline is an **array** used to hold

1. the **stages** to execute
 - each stage is a document (JSON-like)
 2. the **parameters** for that stage
 - the parameters of each stage are stored in the documents
- Conforms to the standard data structure of MongoDB

```
db.cows.aggregate([
  { $match: { ... } },
  { $project: { ... } },
  { $limit: { ... } }
])
```

Suggested ordering of stages in the pipeline

1. `$match` - filter documents early
2. `$project` - reshape documents if needed
3. `$unwind` - flatten arrays if necessary
4. `$group` - perform grouping and aggregations
5. `$sort` - sort the results
6. `$limit` / `$skip` - reduce the result set size
7. `$lookup` - perform joins if needed
8. `$facet` - perform multi-pipeline operations (if needed)
9. `$count` or `$merge` / `$out` - end the pipeline

Stages

\$match

This is similar to a SQL WHERE clause. Used to select documents that match certain conditions

```
// get documents that have farm is equal to 1
{ $match: { "farm": 1 } }
```

```
// We can also use $expr for aggregation expressions. Here we find documents that have a spending greater than its budget
db.monthlybudget.aggregate([
  { $match: {
    $expr: {
```

```

        $gt: ["$spent", "$budget"]
    }
  }
})

```

\$expr

We use `$expr` in match when we need to use calculations or field-to-field comparisons

\$project

Reshapes each document by specifying which fields to include or exclude. In the project stage:

1 means include the field

0 means exclude the field

Unspecified fields are excluded by default. However, `_id` is included by default

```

// Reshape the document by only including the name, milk, but exclude id
{$project: {"name": 1, "milk": 1, "_id": 0}}

```

\$limit

Limits the result set to a specified number. Used when you only need a subset of results

```

// Only want 10
{$limit: 10}

```

\$group

This stage takes the incoming stream of documents, and segments it. Each group is represented by a single document

```

cowcol.aggregate([
  {$group: {
    _id: "$farm",
    total_milk: {$sum: "$milk"}
  }}
])
/*
This results in
{ "_id": 4, "total_milk": 96562 }
{ "_id": 1, "total_milk": 110104 }
{ "_id": 2, "total_milk": 99335 }
{ "_id": 5, "total_milk": 108357 }
{ "_id": 3, "total_milk": 85142 }
*/

```

```

// Find the total marks for each student across all subjects
db.marks.aggregate([
  {$group: {
    _id: "$name",
    "total": {$sum: "$marks"}
  }}
])

```

```

// Find the top two subjects based on average marks
db.mark.aggregate([
  {$group: {
    _id: $name,
    average: {$avg: $marks}
  }},
  {$sort: {average: -1}},
])

```

```
    {$limit: 2}
  })
```

\$addFields

This stage takes the incoming stream of documents, and adds a new field to the document as it is processed

\$function

allows the use of custom Javascript functions within an aggregation pipeline. You must declare the

1. body
2. args
3. lang

```
// Add the age field
db.person.aggregate([
  {$addFields: {
    age: {
      $function: {
        body: function() {
          age = Math.floor((Math.random() * 5) + 1) ;
          return age;
        },
        args: [],
        lang: 'js'
      }
    }
  }}
])
```

\$unwind

Deconstructs an array field outputting a document for each element. It essentially splits a document that contains an array into multiple documents, each with a single value from the array. We need to specify a field path to indicate the array to be deconstructed or specify a document operator.

Options

1. `preserveNullAndEmptyArrays`: setting this option to true preserves documents that have empty arrays
2. `includeArrayIndex`: this option allows you to include the index of each element in the array in the output

Sample Document

```
[
  {
    "_id": 1,
    "name": "Store A",
    "branches": [
      { "locations": ["Downtown", "Uptown"], "employees": 15 },
      { "locations": null, "employees": 8 }
    ]
  },
  {
    "_id": 2,
    "name": "Store B",
    "branches": [
      { "locations": ["Northside"], "employees": 12 },
      { "locations": [], "employees": 20 }
    ]
  },
  {
    "_id": 3,
    "name": "Store C",
    "branches": null
  }
]
```



```
}  
]
```

Unwind operation

```
db.stores.aggregate([  
  {$unwind: {  
    path: "$branches.location",  
    preserveNullAndEmptyArrays: true }  
  }  
])
```

Result

```
{ "_id": 1, "name": "Store A", "branches": { "locations": "Downtown", "employees": 15 } }  
{ "_id": 1, "name": "Store A", "branches": { "locations": "Uptown", "employees": 15 } }  
{ "_id": 1, "name": "Store A", "branches": { "locations": null, "employees": 8 } }  
{ "_id": 2, "name": "Store B", "branches": { "locations": "Northside", "employees": 12 } }  
{ "_id": 2, "name": "Store B", "branches": { "locations": null, "employees": 20 } }  
{ "_id": 3, "name": "Store C", "branches": null }
```

Examples

```
db.budgetitem.aggregate([  
  {  
    $match: {  
      $and: [  
        {  
          $expr: { $gt: ["$spent", "$budget"] } // Filters documents where 'spent' is greater than  
'budget'.  
        },  
        {  
          spent: { $gte: 250 } // Filters documents where 'spent' is greater than or equal to 250.  
        },  
        {  
          "details": { $exists: true } // Ensures the 'details' field exists in the documents.  
        }  
      ]  
    }  
  },  
  {  
    $project: {  
      _id: 0, // Excludes the '_id' field from the output.  
      expense: 1, // Includes the 'expense' field in the output.  
      airline: "$details.airline", // Includes the 'airline' field nested inside 'details'.  
      flight_class: "$details.seat_class", // Includes the 'seat_class' field nested inside 'details' as  
'flight_class'.  
      flight_type: "$details.flight_type", // Includes the 'flight_type' field nested inside 'details'.  
      overspend: { $subtract: ["$spent", "$budget"] } // Calculates the overspend as 'spent - budget'.  
    }  
  }  
])
```

\$lookup

- Performs a left outer join between two collections
- used to combine data from multiple collections based on a specified relationship
- references documents from one collection within another collection

```
{  
  $lookup: {  
    from: "<foreign_collection>"  
    localField: "<field_in_current_collection>",  
    foreignField: "field_in_foreign_collection",
```

```

        as: "output_field",
    }
}

```

from: the name of the other collection you want to join

localField: the field from the current collection that you want to match

foreignField: the field from the collection that you want to match with localField

as: the name of the new array field where the matching documents from the collection will be stored

Orders Collection

```

[
  { "_id": 1, "product": "Pen", "quantity": 10, "customer_id": 101 },
  { "_id": 2, "product": "Notebook", "quantity": 5, "customer_id": 102 },
  { "_id": 3, "product": "Pencil", "quantity": 15, "customer_id": 101 }
]

```

Customers Collection

```

[
  { "_id": 101, "name": "Alice", "location": "New York" },
  { "_id": 102, "name": "Bob", "location": "Los Angeles" },
  { "_id": 103, "name": "Charlie", "location": "Chicago" }
]

```

Query

```

db.orders.aggregate([
  {
    $lookup: {
      from: "customers",
      localField: "customer_id",
      foreignField: "_id",
      as: "customer_details"
    }
  }
])

```

Results

```

[
  {
    "_id": 1,
    "product": "Pen",
    "quantity": 10,
    "customer_id": 101,
    "customer_details": [
      { "_id": 101, "name": "Alice", "location": "New York" }
    ]
  },
  {
    "_id": 2,
    "product": "Notebook",
    "quantity": 5,
    "customer_id": 102,
    "customer_details": [
      { "_id": 102, "name": "Bob", "location": "Los Angeles" }
    ]
  },
  {
    "_id": 3,
    "product": "Pencil",
    "quantity": 15,
    "customer_id": 101,
    "customer_details": [
      { "_id": 101, "name": "Alice", "location": "New York" }
    ]
  }
]

```

```
]
}
]
```

Index

Aggregation Expression: Operators

- `$abs`: Returns the absolute value of a number.
- `$accumulator`: Defines a custom aggregation function in the `$group` stage.
- `$acos`: Returns the arccosine (inverse cosine) of a number in radians.
- `$acosh`: Returns the inverse hyperbolic cosine (arccosh) of a number.
- `$add`: Adds numbers, dates, or arrays.
- `$addToSet`: Adds a value to a set, ensuring no duplicates in the `$group` stage.
- `$allElementsTrue`: Checks if all elements of an array evaluate to true.
- `$and`: Evaluates to true if all expressions evaluate to true.
- `$anyElementTrue`: Checks if any element of an array evaluates to true.
- `$arrayElemAt`: Returns the element at the specified index in an array.
- `$arrayToObject`: Converts an array into a single document.
- `$asin`: Returns the arcsine (inverse sine) of a number in radians.
- `$asinh`: Returns the inverse hyperbolic sine of a number.
- `$atan`: Returns the arctangent (inverse tangent) of a number in radians.
- `$atan2`: Returns the arctangent of two numbers (y and x coordinates).
- `$atanh`: Returns the inverse hyperbolic tangent of a number.
- `$avg`: Calculates the average of numeric values.
- `$binarySize`: Returns the size of a binary value in bytes.
- `$bsonSize`: Returns the size of a BSON document.
- `$ceil`: Rounds a number up to the nearest integer.
- `$cmp`: Compares two values and returns -1, 0, or 1.
- `$concat`: Concatenates strings.
- `$concatArrays`: Combines multiple arrays into a single array.
- `$cond`: Evaluates a conditional expression (if-then-else logic).
- `$convert`: Converts a value to a specified type.
- `$cos`: Returns the cosine of a number in radians.
- `$dateFromParts`: Constructs a date from individual date parts.
- `$dateFromString`: Converts a string to a date object.
- `$dateToParts`: Extracts date parts from a date object.
- `$dateToString`: Formats a date as a string.
- `$dayOfMonth`: Returns the day of the month for a date.
- `$dayOfWeek`: Returns the day of the week for a date.
- `$dayOfYear`: Returns the day of the year for a date.
- `$degreesToRadians`: Converts degrees to radians.
- `$divide`: Divides one number by another.
- `$eq`: Checks if two values are equal.
- `$exp`: Returns e (Euler's number) raised to the power of a specified number.
- `$filter`: Filters elements from an array.
- `$first`: Returns the first element in an array.
- `$floor`: Rounds a number down to the nearest integer.
- `$function`: Defines a custom function to execute in aggregation.
- `$gt`: Checks if a value is greater than another value.

- `$gte`: Checks if a value is greater than or equal to another value.
- `$hour`: Extracts the hour from a date object.
- `$ifNull`: Returns a specified value if the input is null or missing.
- `$in`: Checks if a value is in an array.
- `$indexOfArray`: Returns the index of an element in an array.
- `$indexOfBytes`: Returns the byte index of a substring in a string.
- `$indexOfCP`: Returns the code point index of a substring in a string.
- `$isArray`: Checks if the input is an array.
- `$isNumber`: Checks if the input is a number.
- `$isoDayOfWeek`: Returns the ISO day of the week for a date.
- `$isoWeek`: Returns the ISO week of the year for a date.
- `$isoWeekYear`: Returns the ISO week-numbering year for a date.
- `$last`: Returns the last element in an array.
- `$let`: Assigns variables for use in an aggregation pipeline.
- `$literal`: Returns a value without parsing or interpretation.
- `$ln`: Returns the natural logarithm of a number.
- `$log`: Returns the logarithm of a number for a specified base.
- `$log10`: Returns the base-10 logarithm of a number.
- `$lt`: Checks if a value is less than another value.
- `$lte`: Checks if a value is less than or equal to another value.
- `$ltrim`: Removes leading whitespace or characters from a string.
- `$map`: Applies an expression to each element in an array.
- `$max`: Returns the maximum value in a set of values.
- `$mergeObjects`: Combines multiple documents into a single document.
- `$meta`: Accesses metadata in a query.
- `$millisecond`: Extracts the millisecond from a date object.
- `$min`: Returns the minimum value in a set of values.
- `$minute`: Extracts the minute from a date object.
- `$mod`: Returns the remainder of a division operation.
- `$multiply`: Multiplies two or more numbers.
- `$ne`: Checks if two values are not equal.
- `$not`: Negates a specified expression.
- `$objectToArray`: Converts a document into an array of key-value pairs.
- `$or`: Evaluates to true if any expression evaluates to true.
- `$pow`: Raises a number to a specified power.
- `$push`: Appends a value to an array in the `$group` stage.
- `$radiansToDegrees`: Converts radians to degrees.
- `$range`: Generates an array of sequential numbers.
- `$reduce`: Applies an expression to reduce an array to a single value.
- `$regexFind`: Finds a match for a regular expression in a string.
- `$regexFindAll`: Finds all matches for a regular expression in a string.
- `$regexMatch`: Checks if a string matches a regular expression.
- `$replaceOne`: Replaces the first match of a string or regex with a specified value.
- `$replaceAll`: Replaces all matches of a string or regex with a specified value.
- `$reverseArray`: Reverses the order of elements in an array.
- `$round`: Rounds a number to a specified decimal place.
- `$rtrim`: Removes trailing whitespace or characters from a string.
- `$second`: Extracts the second from a date object.
- `$setDifference`: Returns elements in one array but not in another.

- `$setEquals`: Checks if two arrays contain the same elements.
- `$setIntersection`: Returns elements common to all input arrays.
- `$setIsSubset`: Checks if one array is a subset of another array.
- `$setUnion`: Combines elements from all input arrays without duplicates.
- `$sin`: Returns the sine of a number in radians.
- `$size`: Returns the size of an array.
- `$slice`: Extracts a subset of an array.
- `$split`: Splits a string into an array of substrings.
- `$sqrt`: Returns the square root of a number.
- `$stdDevPop`: Calculates the population standard deviation of numeric values.
- `$stdDevSamp`: Calculates the sample standard deviation of numeric values.
- `$strLenBytes`: Returns the number of bytes in a string.
- `$strLenCP`: Returns the number of UTF-8 code points in a string.
- `$strcasecmp`: Performs a case-insensitive string comparison.
- `$substr`: Extracts a substring from a string.
- `$substrBytes`: Extracts a substring based on byte indexes.
- `$substrCP`: Extracts a substring based on UTF-8 code point indexes.
- `$sum`: Calculates the sum of numeric values.
- `$switch`: Evaluates multiple conditions and returns a value for the first true condition.
- `$tan`: Returns the tangent of a number in radians.
- `$toBool`: Converts a value to a boolean.
- `$toDate`: Converts a value to a date.
- `$toDecimal`: Converts a value to a decimal.
- `$toDouble`: Converts a value to a double.
- `$toInt`: Converts a value to an integer.
- `$toLong`: Converts a value to a 64-bit integer.
- `$toLower`: Converts a string to lowercase.
- `$toObjectId`: Converts a value to an ObjectId.
- `$toString`: Converts a value to a string.
- `$toUpper`: Converts a string to uppercase.
- `$trim`: Removes leading and trailing whitespace or characters from a string.
- `$trunc`: Truncates a number to a whole number or specified decimal place.
- `$type`: Returns the BSON type of a field.
- `$week`: Returns the week of the year for a date.
- `$year`: Returns the year of a date.
- `$zip`: Combines arrays into a single array of subarrays.